

# ATWONGERE VIANNEY

## 2023/U/MMU/BSSE/00574.

### LOGISTIC REGRESSION AND FINE TUNING IT

```
In [1]: 1 #Required Libraries
        2 from sklearn.datasets import make_classification
        3 from sklearn.model_selection import GridSearchCV
        4 from sklearn.linear_model import LogisticRegression
        5 from sklearn.model_selection import train_test_split
        6 import pandas as pd
```

```
In [2]: 1 df=pd.read_csv("C:\\Users\\Atwongire Vianney\\Desktop\\AI_PRAC\\03_Cluster
        2 df
```

Out[2]:

	gradyear	gender	age	NumberOffriends	basketball	football	soccer	softball	volleyba
0	2007	NaN	NaN	0	0	0	0	0	
1	2007	F	17.41	49	0	0	1	0	
2	2007	F	17.511	41	0	0	0	0	
3	2006	F	NaN	36	0	0	0	0	
4	2008	F	16.657	1	0	0	0	0	
...	...	...	...	...	...	...	...	...	...
14995	2008	F	16.329	21	0	0	0	0	
14996	2008	F	16.545	50	0	0	0	0	
14997	2007	F	17.999	32	0	0	0	0	
14998	2007	F	17.903	20	0	0	0	0	
14999	2009	F	15.811	25	0	0	7	0	

15000 rows × 40 columns



```
1 df.head()
```

	gradyear	gender	age	NumberOffriends	basketball	football	soccer	softball	volleyball	swimming
0	2007	NaN	NaN	0	0	0	0	0	0	0
1	2007	F	17.41	49	0	0	1	0	0	0
2	2007	F	17.511	41	0	0	0	0	0	0
3	2006	F	NaN	36	0	0	0	0	0	0
4	2008	F	16.657	1	0	0	0	0	0	0

◀ ▶

```
1 df.tail()
```

	gradyear	gender	age	NumberOffriends	basketball	football	soccer	softball	volleyba
14995	2008	F	16.329	21	0	0	0	0	
14996	2008	F	16.545	50	0	0	0	0	
14997	2007	F	17.999	32	0	0	0	0	
14998	2007	F	17.903	20	0	0	0	0	
14999	2009	F	15.811	25	0	0	7	0	

◀ ▶

```
1 x=df['NumberOfFriends'].array.reshape(-1,1)
2 x
```

```
[
[0],
[49],
[41],
[36],
[1],
[32],
[18],
[0],
[0],
[21],
[0],
[0],
[29],
[0],
[89],
[37],
[89],
[22]
```

```
In [6]: 1 y=df['sex']
        2 y
```

```
Out[6]: 0      0
        1      0
        2      3
        3      0
        4      0
        ..
       14995    0
       14996    0
       14997    0
       14998    0
       14999    0
       Name: sex, Length: 15000, dtype: int64
```

```
In [7]: 1 #splitting the data
        2 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_st
```

```
In [8]: 1 x_train.shape
```

```
Out[8]: (12000, 1)
```

```
In [9]: 1 model=LogisticRegression()
        2 model.fit(x_train,y_train)
```

C:\Users\Atwongire Vianney\anaconda3\Lib\site-packages\sklearn\linear\_model\\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)  
 Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))  
 n\_iter\_i = \_check\_optimize\_result(

```
Out[9]: ▾ LogisticRegression
        LogisticRegression()
```

```
In [10]: 1 y_pred=model.predict(x_test)
         2 y_pred
```

```
Out[10]: array([0, 0, 0, ..., 0, 0, 0], dtype=int64)
```

```
In [11]: 1 score=model.score(x_train,y_train)
         2 score
```

```
Out[11]: 0.89125
```

## FINE TUNING

```
In [12]: 1 #Defining parameters
2 param_grid={
3     'penalty':['l1','l2'],
4     'C':[0.001,0.01,0.1,1,10,100],
5     'solver':['liblinear','saga']
6 }
```

```
In [13]: 1 #performing grid with cross_validation
2 grid_search=GridSearchCV(estimator=model,param_grid=param_grid,cv=5,n_jobs=
3 grid_search.fit(x_train,y_train)
```

C:\Users\Atwongere Vianney\anaconda3\Lib\site-packages\sklearn\model\_selection\\_split.py:725: UserWarning: The least populated class in y has only 1 members, which is less than n\_splits=5.  
warnings.warn(

```
Out[13]:  ▸ GridSearchCV
  ▸ estimator: LogisticRegression
    ▸ LogisticRegression
```

```
In [14]: 1 #best parameters and best score
2 best_param=grid_search.best_params_
3 grid_search
```

```
Out[14]:  ▸ GridSearchCV
  ▸ estimator: LogisticRegression
    ▸ LogisticRegression
```

```
In [15]: 1 best_score=grid_search.best_score_
2 best_score
```

```
Out[15]: 0.8914166666666666
```

```
In [16]: 1 #evaluating the performance
2 score=grid_search.score(x_test,y_test)
3 score
```

```
Out[16]: 0.9046666666666666
```

```
In [17]: 1 y_pred=model.predict(x_test)
2 y_pred
```

```
Out[17]: array([0, 0, 0, ..., 0, 0, 0], dtype=int64)
```

## FOR LINEAR REGRESSION

```
In [18]: 1 import pandas as pd
2 from sklearn.linear_model import LinearRegression
3 from sklearn.model_selection import train_test_split
```

```
In [19]: 1 df=pd.read_csv("C:\\Users\\Atwongere Vianney\\Desktop\\AI_PRAC\\03_Cluster")
2 df
```

Out[19]:

	gradyear	gender	age	NumberOfriends	basketball	football	soccer	softball	volleyba
0	2007	NaN	NaN	0	0	0	0	0	
1	2007	F	17.41	49	0	0	1	0	
2	2007	F	17.511	41	0	0	0	0	
3	2006	F	NaN	36	0	0	0	0	
4	2008	F	16.657	1	0	0	0	0	
...	...	...	...	...	...	...	...	...	...
14995	2008	F	16.329	21	0	0	0	0	
14996	2008	F	16.545	50	0	0	0	0	
14997	2007	F	17.999	32	0	0	0	0	
14998	2007	F	17.903	20	0	0	0	0	
14999	2009	F	15.811	25	0	0	7	0	

15000 rows × 40 columns



```
In [20]: 1 df.head()
```

Out[20]:

	gradyear	gender	age	NumberOfriends	basketball	football	soccer	softball	volleyball	st
0	2007	NaN	NaN	0	0	0	0	0	0	
1	2007	F	17.41	49	0	0	1	0	0	
2	2007	F	17.511	41	0	0	0	0	0	
3	2006	F	NaN	36	0	0	0	0	0	
4	2008	F	16.657	1	0	0	0	0	0	

5 rows × 40 columns





```
In [24]: 1 #splitting the data
          2 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_st
```

```
In [25]: 1 model=LinearRegression()
          2 model.fit(x_train,y_train)
```

```
Out[25]: ▾ LinearRegression
          LinearRegression()
```

```
In [26]: 1 score=model.score(x_train,y_train)
          2 score
```

```
Out[26]: 9.581120157475809e-06
```

```
In [27]: 1 model.coef_
```

```
Out[27]: array([0.00011896])
```

```
In [28]: 1 model.intercept_
```

```
Out[28]: 0.2200046732690924
```

## FINE TUNING

```
In [29]: 1 #Defining parameters
          2 from sklearn.model_selection import GridSearchCV
          3 from sklearn.metrics import mean_squared_error
```

```
In [30]: 1 param_grid={
          2     'copy_X':[True,False],
          3     'fit_intercept':[True,False],
          4     'n_jobs':[True,False],
          5     'positive':[True,False]
          6
          7 }
```

```
In [31]: 1 #performing grid with cross_validation
          2 grid_search=GridSearchCV(model,param_grid,cv=5,scoring='neg_mean_squared_e
          3 grid_search.fit(x_train,y_train)
```

```
Out[31]: ▸ GridSearchCV
          ▸ estimator: LinearRegression
              ▸ LinearRegression
```

```
In [32]: 1 #Best model
          2 best_model=grid_search.best_estimator_
          3 best_model
```

```
Out[32]: LinearRegression
          LinearRegression(n_jobs=True, positive=True)
```

```
In [33]: 1 best_score=model.score(x_test,y_test)
          2 best_score
```

```
Out[33]: -0.0035208057122375624
```

```
In [34]: 1 y_pred=best_model.predict(x_test)
          2 y_pred
```

```
Out[34]: array([0.22143224, 0.2202426 , 0.22143224, ..., 0.22464428, 0.2202426 ,
                0.22285981])
```

```
In [35]: 1 mse=mean_squared_error(y_test,y_pred)
          2 mse
```

```
Out[35]: 0.6336346329671283
```

```
In [ ]: 1
```